

# BASANT BHATTARAI

## SENIOR DATA & AI ENGINEER / PLATFORMS, DATABASES & AGENTIC SYSTEMS

Kathmandu, Nepal | [technobasant9@gmail.com](mailto:technobasant9@gmail.com) | [linkedin.com/in/technobasant](https://linkedin.com/in/technobasant) | [github.com/technobasant](https://github.com/technobasant)

### PROFILE

Senior data and AI engineer with 9 years in software engineering, six of them owning production data platforms. Open-source first — Kafka, Spark, Iceberg, Trino, dbt and Airflow on Kubernetes — across ingestion, transformation and analytical serving, plus the databases underneath: modeling, replication, backup and recovery, upgrades. Designs and runs the agent workflows on that same stack — LangGraph, Google ADK and MCP tool servers — so model output lands as typed, reviewable rows rather than a transcript. Track record: 40% lower infrastructure cost, 50% faster processing, 60% faster analytics delivery.

### CORE STRENGTHS

|                                                                                                                 |                                                                                                                |
|-----------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|
| <b>Platform ownership</b><br>Consumer-facing reliability, governance, recovery, and operational readiness.      | <b>Data engineering</b><br>Batch and streaming pipelines, orchestration, modeling, and change-safe processing. |
| <b>Database lifecycle</b><br>Modeling, replication, failover, backup and PITR, upgrades, and query performance. | <b>Agentic systems</b><br>LangGraph, Google ADK and MCP tool servers, evaluation, and typed persistence.       |
| <b>Technical leadership</b><br>Design review, mentoring, runbooks, and making system ownership transferable.    | <b>Distributed work</b><br>Long-term collaboration across European and US time zones from Kathmandu.           |

### EXPERIENCE

#### Senior Data Engineer | UXCam

Feb 2024 - present

- Own reliability and governance for a platform handling 15 TB+ and 100M+ events daily: freshness and latency objectives per dataset, 99.9% uptime, incident-ready runbooks, reviewable data contracts, and a named owner for every published table.
- Design and operate the production agent workflows — LangGraph and Google ADK orchestration, MCP tool servers, retrieval and evaluation — against the existing platform, so model output lands as typed, provenanced rows rather than a transcript: analytics delivery 60% faster with 75% less analyst effort, review and fallback paths kept.
- Lead architecture and operational review across ingestion, transformation, storage, analytical serving, observability and recovery, without treating any single tool as the architecture.
- Run capacity and cost work as a design input rather than a monthly report, using storage layout, partitioning and compaction cadence as the primary levers.
- Mentor 5 engineers through pairing, code review and written design feedback; team delivery time down 30%, with the emphasis on making system ownership transferable.

#### Data Engineer | UXCam

Feb 2020 - Feb 2024

- Built and operated batch and streaming pipelines, then took on the data models, orchestration, storage and analytical serving paths around them.
- Rebuilt the core ETL layer in Spark and PySpark with Airflow orchestration, cutting processing time for 5 TB+/day by 50% and unblocking real-time product dashboards.
- Made change and recovery safe by treating schema evolution, replay, late-arriving data and backfills as designed interfaces with idempotent loads and explicit watermarks, not as emergency procedures.

- 
- Tuned Trino, ClickHouse, Citus and TimescaleDB serving 100M+ queries/day for a 50% p95 latency reduction through partitioning, clustering and model redesign.
  - Led the Kubernetes migration of the data infrastructure: 40% lower infrastructure cost, 99.9% uptime, and autoscaling that held under peak load.
  - Worked across European and US time zones to turn product questions into maintainable datasets.

#### **Project Leader | SVCET, India**

2019 - 2020

- Led 4 developers building an intelligent dialogue system on an NLP and deep-learning stack, reaching the client's acceptance threshold for intent recognition and delivering 2 weeks ahead of schedule.
- Coordinated requirements, planning and review across academic and industry stakeholders, and introduced the code-review and version-control discipline the department kept afterwards.

#### **Software Developer | SV Technology, India**

2017 - 2019

- Built Python and FastAPI backend services carrying 10K+ users at 85% test coverage.
- Designed and maintained PostgreSQL schemas, stored procedures and reporting paths, and set up CI with Jenkins alongside the early pipeline work that moved me into data engineering.

### **CAPABILITY MAP**

#### **Platform ownership**

Python and SQL; Kafka ingestion; Spark/PySpark processing; Airflow orchestration; Iceberg, dbt and lakehouse modeling; Kubernetes, Docker and Terraform delivery; contracts, lineage, cost and recovery.

#### **Database lifecycle**

PostgreSQL, ClickHouse, Trino, Citus, TimescaleDB, Redis and MongoDB; transactional and analytical modeling, partitioning, replication, backup, restore, high availability and query performance.

#### **Agentic systems**

LangGraph, LangChain, Google ADK, MCP servers, Pydantic, Milvus and MLflow; bounded tools, retrieval, evaluation before rollout, typed output, provenance, human review and honest fallback paths.

#### **Technical leadership**

Architecture and operational review, mentoring, incident-ready runbooks, service and freshness objectives, observability with Prometheus and Grafana, and distributed delivery across time zones.

### **EDUCATION AND LANGUAGES**

#### **Bachelor of Technology in Computer Science and Engineering**

JNTUA College of Engineering, Anantapur, India | 2015 - 2019

English: C1 | Nepali: native

**Selected work, writing and case studies:** [technobasant.github.io](https://technobasant.github.io)