

Basant Bhattarai

Senior Data & AI Engineer · Platform & Database Reliability

technobasant9@gmail.com · linkedin.com/in/technobasant · github.com/technobasant · Working US and EU hours

Experience at a glance

Role	Period	Public scope
Senior Data Engineer, UXCam	2024 – present	Data-platform and AI-product reliability, governance, architecture review, and technical leadership
Data Engineer, UXCam	2020 – 2024	Batch and streaming pipelines, data modeling, orchestration, storage, serving, and operations
Project Leader, SVCET	2019 – 2020	Planning and delivery of an NLP dialogue-system project with a small team
Software Developer, SV Technology	2017 – 2019	Python backend services, PostgreSQL schema work, reporting, and CI

Senior Data Engineer · UXCam

February 2024 – present · Working with a global team

- Own reliability and governance outcomes across production data and AI-assisted workflows, including clear service expectations, incident-ready runbooks, and reviewable data contracts.
- Design boundaries that turn model output into typed, validated, traceable records before another product workflow can depend on it.
- Lead design reviews and operational improvements across data processing, storage, serving, observability, and recovery without treating any one tool as the architecture.
- Mentor engineers through pairing and code review, with an emphasis on making ownership transferable rather than concentrating system knowledge in one person.

Data Engineer · UXCam

February 2020 – February 2024

- Built and operated batch and streaming data pipelines, then took on the data models, orchestration, storage, and analytical serving paths around them.
- Improved recovery and change safety by treating schema evolution, replay, late data, and backfills as designed interfaces instead of emergency procedures.
- Tuned distributed processing and query workloads by starting with the read path, measurement, and data layout before changing infrastructure settings.
- Worked with teams across time zones to turn product questions into maintainable datasets and operational workflows.

Skills

Sustained production experience

Technologies I have operated, debugged, and reviewed as part of long-lived systems.

Technology	Since	Public description	Evidence
Python	2017	Backend services, data pipelines, automation, and AI tooling	Career summary
SQL	2017	Transactional, analytical, and reliability work	PostgreSQL tutorial
Apache Spark / PySpark	2020	Batch and streaming data processing	Career summary
Apache Kafka	2020	Event-driven and streaming systems	Career summary
Apache Airflow	2020	Workflow orchestration, replay, and recovery design	Career summary
Apache Iceberg	2022	Table evolution, maintenance, and data modeling	Career summary
dbt	2021	Transformation contracts, testing, and documentation	Career summary
PostgreSQL	2017	Application data, replication, recovery, and upgrades	PostgreSQL tutorial
ClickHouse	2021	Analytical modeling and operational practice	Career summary
Trino	2021	Federated SQL and query modeling	Career summary
Kubernetes	2022	Container orchestration and data workloads	Career summary
Prometheus + Grafana	2021	Service and data-reliability signals	Reliability lab

Shipped on a narrower surface

Used to deliver real systems, with less breadth or duration than Tier 1.

AI & agents LangGraph · LangChain · CrewAI · Google-ADK · MCP · Milvus · MLflow · Pydantic ·

Data platform Databricks + Unity Catalog · Amazon Kinesis · Spark Structured Streaming · Debezium / CDC · Delta Lake ·

Databases Redis · MongoDB · Citus · TimescaleDB ·

Product & infra FastAPI · Celery · Next.js · Terraform · Docker · GitHub Actions · Helm ·

Clouds AWS: S3, EKS, Lambda, Glue, Athena, Redshift, Kinesis, EMR, RDS, Step Functions, Bedrock, SageMaker · GCP: BigQuery, Dataflow, Dataproc, Datastream, Pub/Sub, Cloud Composer, GKE, Vertex AI · Azure: Synapse, Data Factory, Databricks, ADLS, Event Hubs, AKS, Azure ML ·

Working knowledge

Evaluated, prototyped, or run in a lab; listed to show where the depth stops.

Apache Flink, Apache Beam, Prefect, LlamaIndex, AutoGen, Haystack, Semantic Kernel, DSPy, OpenAI Agents SDK, Pinecone, pgvector, Chroma, FAISS, Weaviate, ScyllaDB, MariaDB + Galera, SolrCloud, Elasticsearch, Apache Druid, Snowflake, BigQuery, Amazon Redshift, DuckDB, Hadoop, Apache Hive, Scala, Java, Azure Synapse, Azure Data Factory, Power BI, Tableau, Looker, Superset, Metabase, PyTorch, Polars, PyArrow, Ansible, Jenkins, LangSmith, Langfuse, Opik, OpenTelemetry, Datadog.

Education and languages

Bachelor of Technology, Computer Science and Engineering — JNTUA College of Engineering, Anantapur, India, 2015–2019.

English: C1 · Nepali: native

Contact

technobasant9@gmail.com · LinkedIn · GitHub